

Failure Detection and Membership II




ECE 599 / CS 519 – SPRING 2015

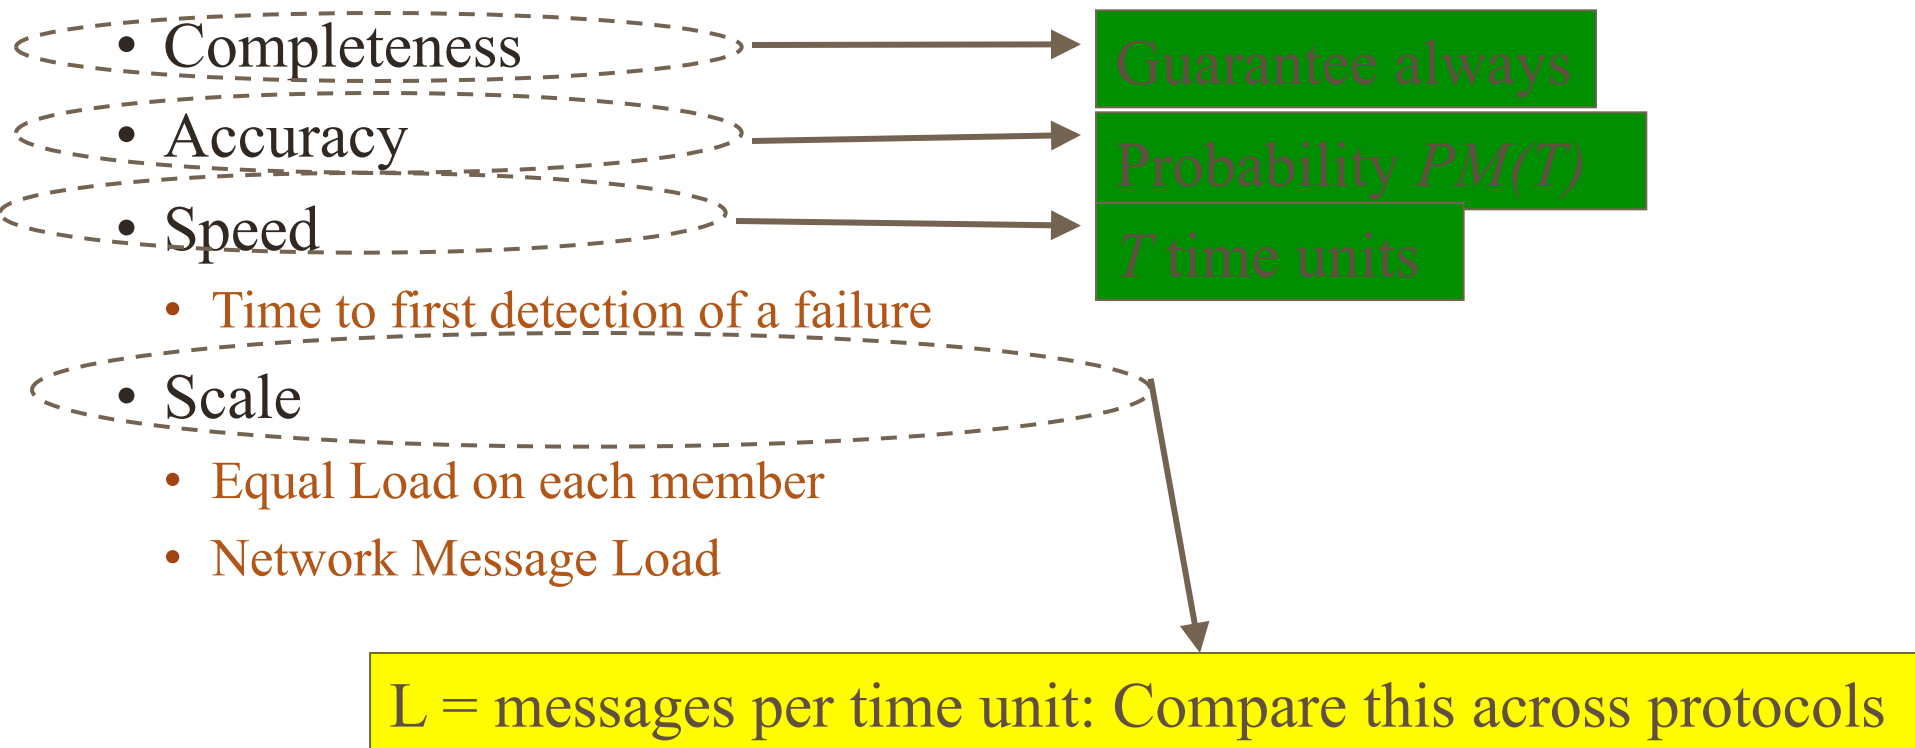
Failure Detector Properties ...

- Completeness
- Accuracy
- Speed
 - Time to first detection of a failure
- Scale
 - Equal Load on each member
 - Network Message Load


...Are application-defined Requirements



...Are application-defined Requirements




All-to-All Heartbeating



Gossip-style Heartbeating

Array of
Heartbeat Seq. l
for member subset

Every t_g units
=gossip period,
send $O(N)$ gossip
message



What's the Best/Optimal we can do?

- *Worst case* load L^* **per member** in the group (messages per time unit)
 - as a function of $T, PM(T), N$
 - Independent Message Loss probability p_{ml}

$$L^* = \frac{\log(PM(T))}{\log(p_{ml})} \cdot \frac{1}{T}$$

Heartbeating

- Optimal L is independent of N (!)
- All-to-all and gossip-based: sub-optimal
 - $L=O(N/T)$
 - try to achieve simultaneous detection at *all* processes
 - fail to distinguish *Failure Detection* and *Dissemination* components


⇒ Key:

- Separate the two components
- Use a non heartbeat-based Failure Detection Component

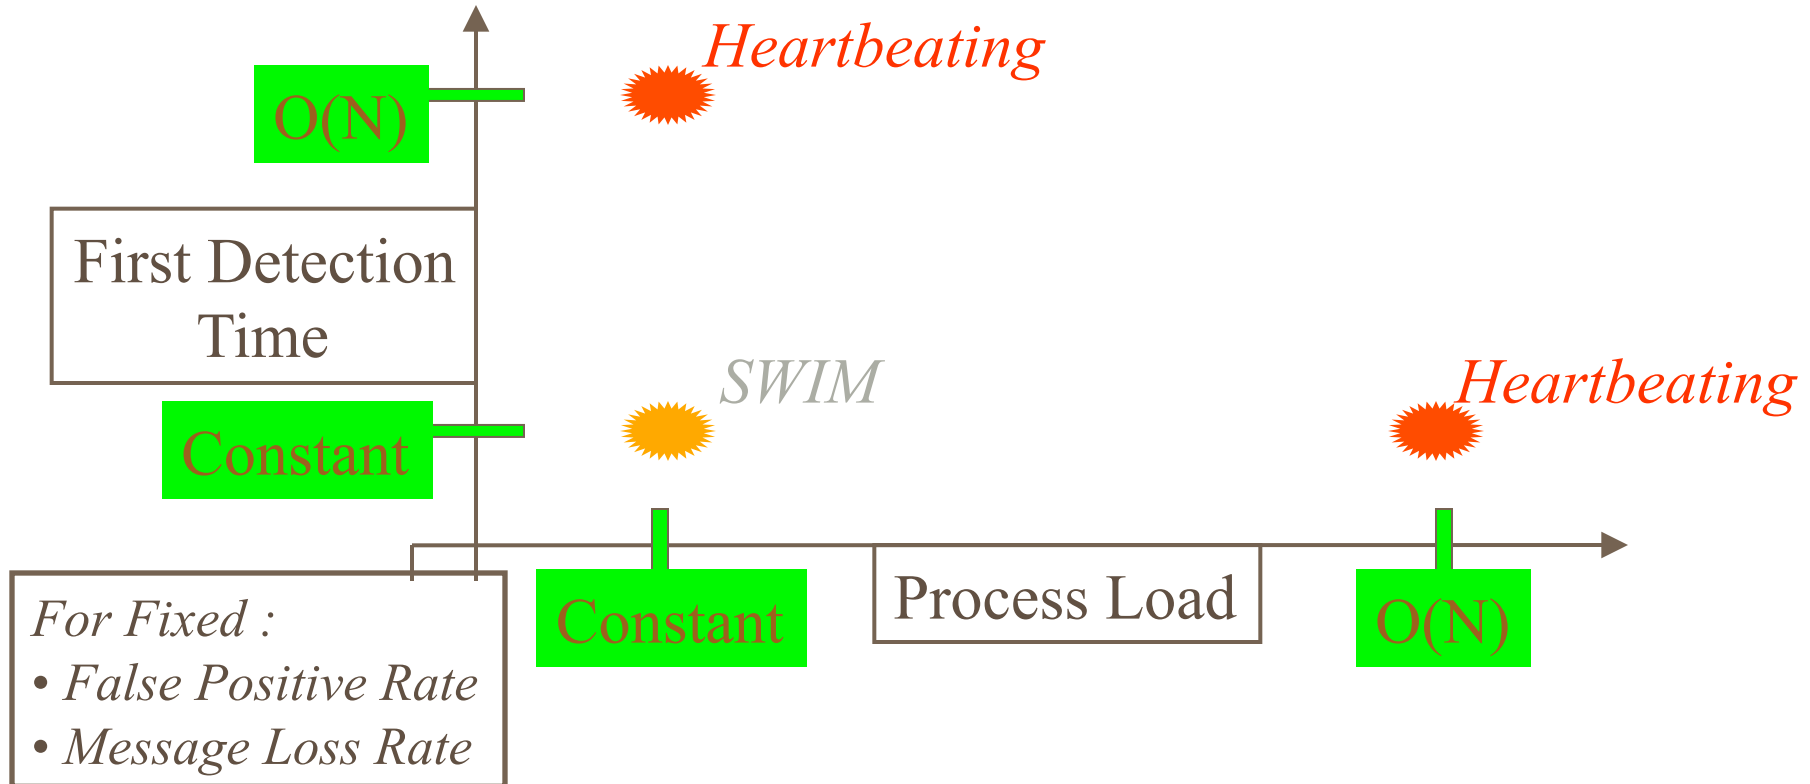
Next

- Is there a better failure detector?

SWIM Failure Detector Protocol



SWIM versus Heartbeating



SWIM Failure Detector

Parameter	SWIM
First Detection Time	<ul style="list-style-type: none">Expected $\left[\frac{e}{e-1} \right]$ periodsConstant (independent of group size)
Process Load	<ul style="list-style-type: none">Constant per period$< 8 L^*$ for 15% loss
False Positive Rate	<ul style="list-style-type: none">Tunable (via K)Falls exponentially as load is scaled
Completeness	<ul style="list-style-type: none">Deterministic time-boundedWithin $O(\log(N))$ periods w.h.p.

Accuracy, Load

$PM(T)$ is exponential in $-K$. Also depends on pml (and pf)

- See paper

$$\frac{L}{L^*} < 28$$

$$\frac{E[L]}{L^*} < 8$$

for up to 15 % loss rates

Detection Time

$$\text{Prob. of being pinged in } T' = 1 - \left(1 - \frac{1}{N}\right)^{N-1} = 1 - e^{-1}$$

$$E[T] = T' \cdot \frac{e}{e-1}$$

Completeness: *Any* alive member detects failure

- Eventually
- By using a trick: within worst case $O(N)$ protocol periods

Next

- How do failure detectors fit into the big picture of a group membership protocol?
- What are the missing blocks?

Group Membership Protocol II

Failure Detector

Some process
finds out quickly

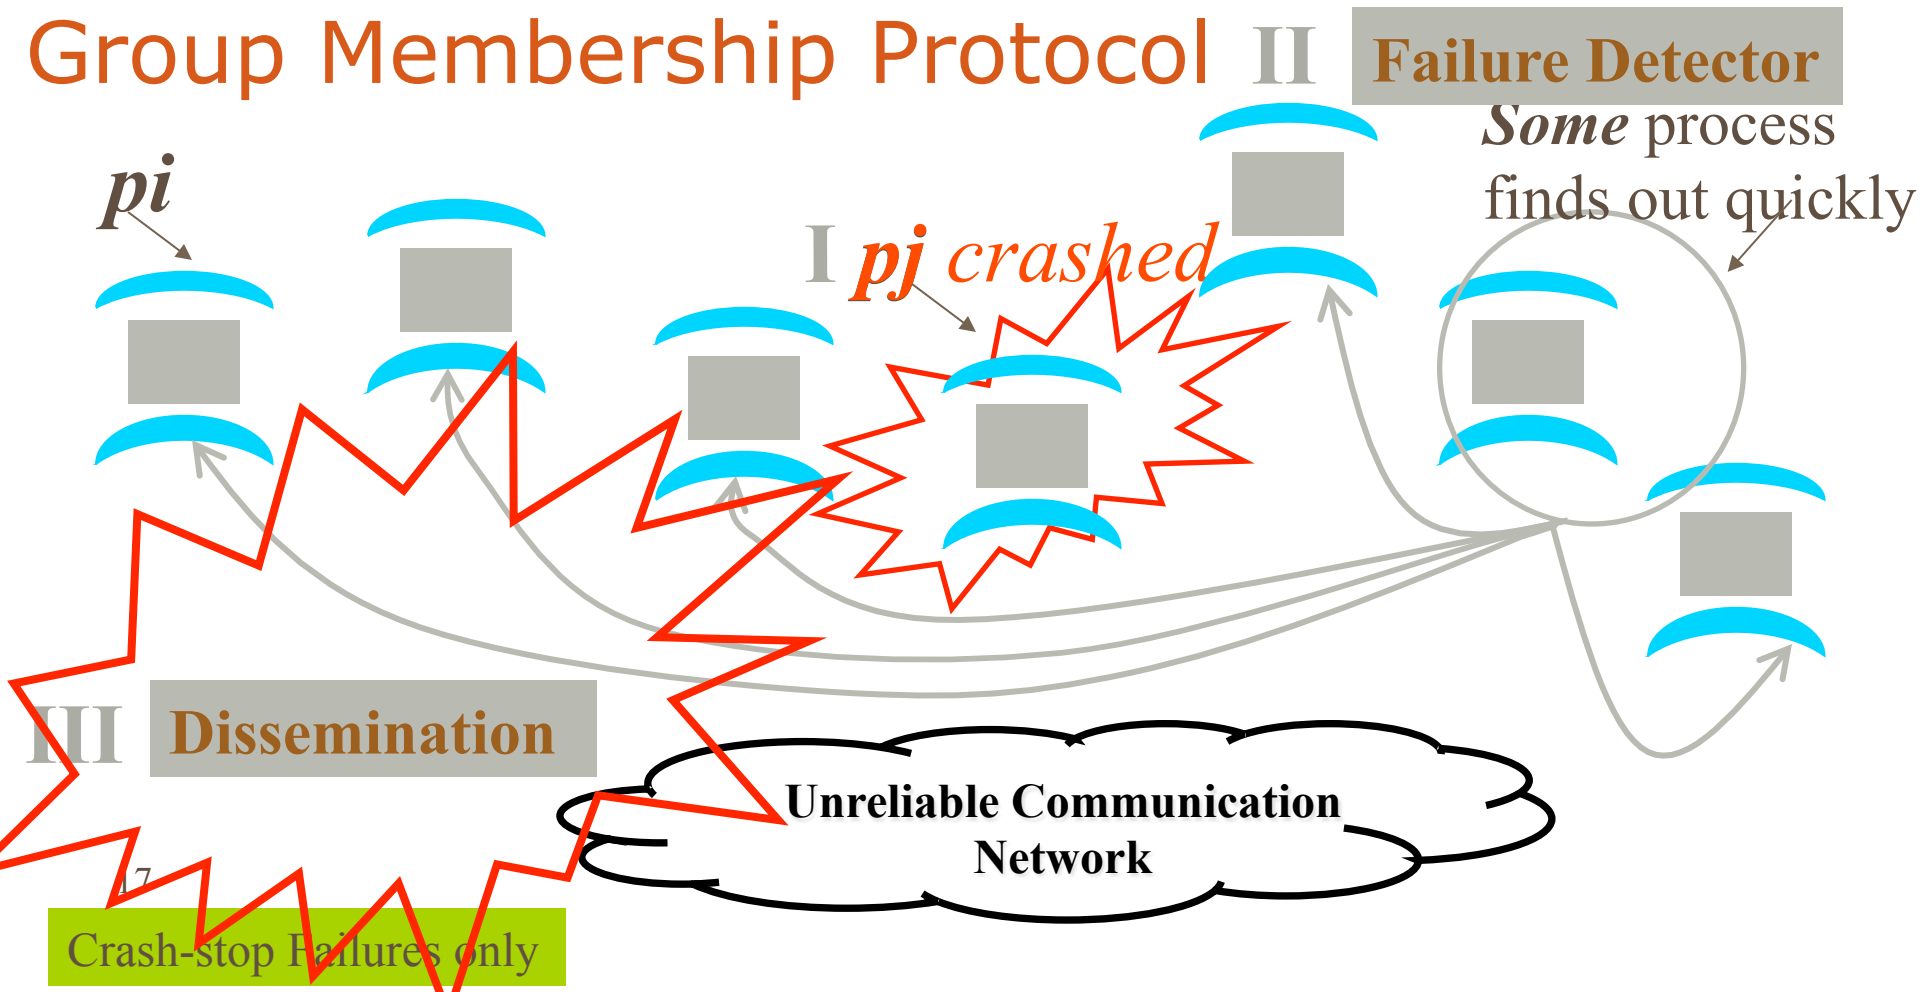
I *p_j crashed*

p_i

Dissemination

Unreliable Communication
Network


Crash-stop Failures only



Dissemination Options

- Multicast (Hardware / IP)
 - unreliable
 - multiple simultaneous multicasts
- Point-to-point (TCP / UDP)
 - expensive
- Zero extra messages: Piggyback on Failure Detector messages
 - Infection-style Dissemination

Infection-style Dissemination



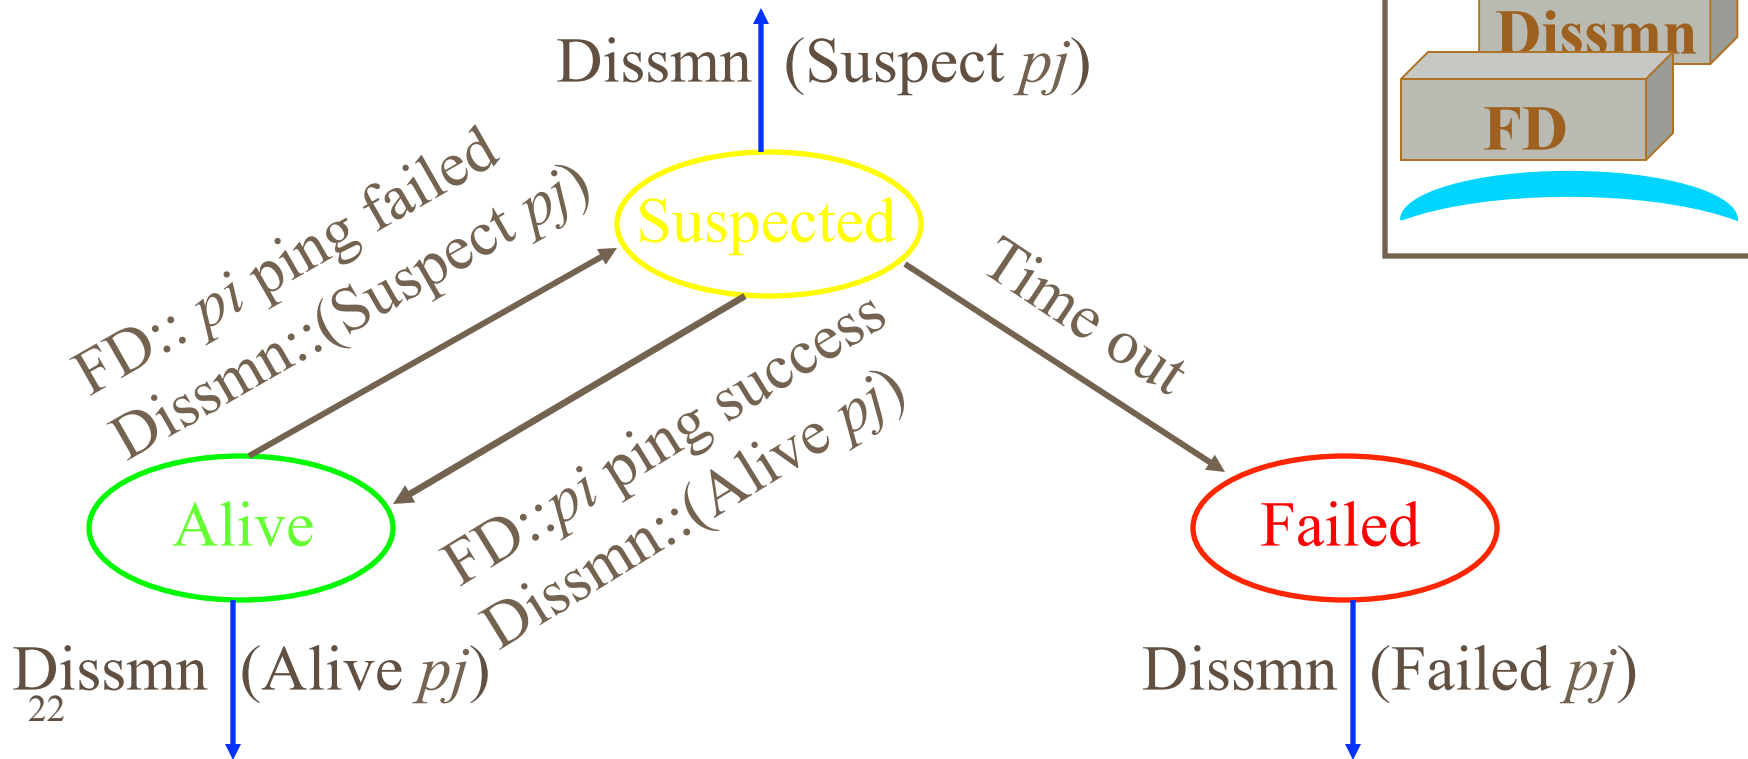
Infection-style Dissemination

- Epidemic/Gossip style dissemination
 - After $\lambda \cdot \log(N)$ protocol periods, $N^{-(2\lambda-2)}$ processes would not have heard about an update
- Maintain a buffer of recently joined/evicted processes
 - Piggyback from this buffer
 - Prefer recent updates
- Buffer elements are garbage collected after a while
 - After $\lambda \cdot \log(N)$ protocol periods, i.e., once they've propagated through the system; this defines weak consistency

Suspicion Mechanism

- False detections, due to
 - Perturbed processes
 - Packet losses, e.g., from congestion
- Indirect pinging may not solve the problem
- Key: *suspect* a process before *declaring* it as failed in the group

Suspicion Mechanism



Suspicion Mechanism

- Distinguish multiple suspicions of a process
 - Per-process *incarnation number*
 - *Inc #* for pi can be incremented only by pi
 - e.g., when it receives a (Suspect, pi) message
 - Somewhat similar to DSDV
- Higher inc# notifications over-ride lower inc#'s
- Within an inc#: (Suspect inc #) > (Alive, inc #)
- (Failed, inc #) overrides everything else

Wrap Up

- Failures the norm, not the exception in datacenters
- Every distributed system uses a failure detector
- Many distributed systems use a membership service
- Ring failure detection underlies
 - IBM SP2 and many other similar clusters/machines
- Gossip-style failure detection underlies
 - Amazon EC2/S3 (rumored!)